

Lecture 11: SQL Join

BADM/ACCY 352

Spring 2017

Instructor: Yi Yang, PhD

Class project announcement

- Previous projects:
 - Carle Hospital
 - MTD bus company
 - Uber
 - Accounting information system
 - Twitter
 - Restaurant
 - etc

Previous lectures

- Data retrieval with SELECT... FROM... command
- However, the real power of SQL is yet to come

This lecture

- We will learn SQL JOIN, the real power of relational database.

Motivation

Gallery

GAL_NUM	GAL_OWNER	GAL_PHONE
5	Alice	123-4456
6	Waters	353-2243

Painter

PTR_NUM	PTR_NAME	PTR_AREACODE	PTR_PHONE
123	Ross	901	885-4567
126	Itero	901	346-1112
127	Geoff	615	221-4456

Painting

PTNG_NUM	PTNG_TITLE	PTNG_PRICE	PTR_NUM	GAL_NUM
1338	Dawn Thunder	245.50	123	5
1339	A Faded Rose	6723.00	123	NULL
1340	The Founders	567.99	126	6
1341	Hasty Pudding Exit	145.50	123	NULL
1342	Plastic Paradise	8328.99	126	6
1343	Roamin'	785.00	127	6
1344	Wild Waters	999.00	127	5
1345	Stuff 'n Such 'n Some	9800.00	123	5

Motivation (cont.)

- List the information of painter who creates the painting 'Dawn Thunder'?

```
SELECT PTR_NUM FROM PAINTING WHERE PTNG_TITLE = 'Dawn Thunder';  
# return 123  
SELECT * FROM PAINTER WHERE PTR_NUM = 123;
```

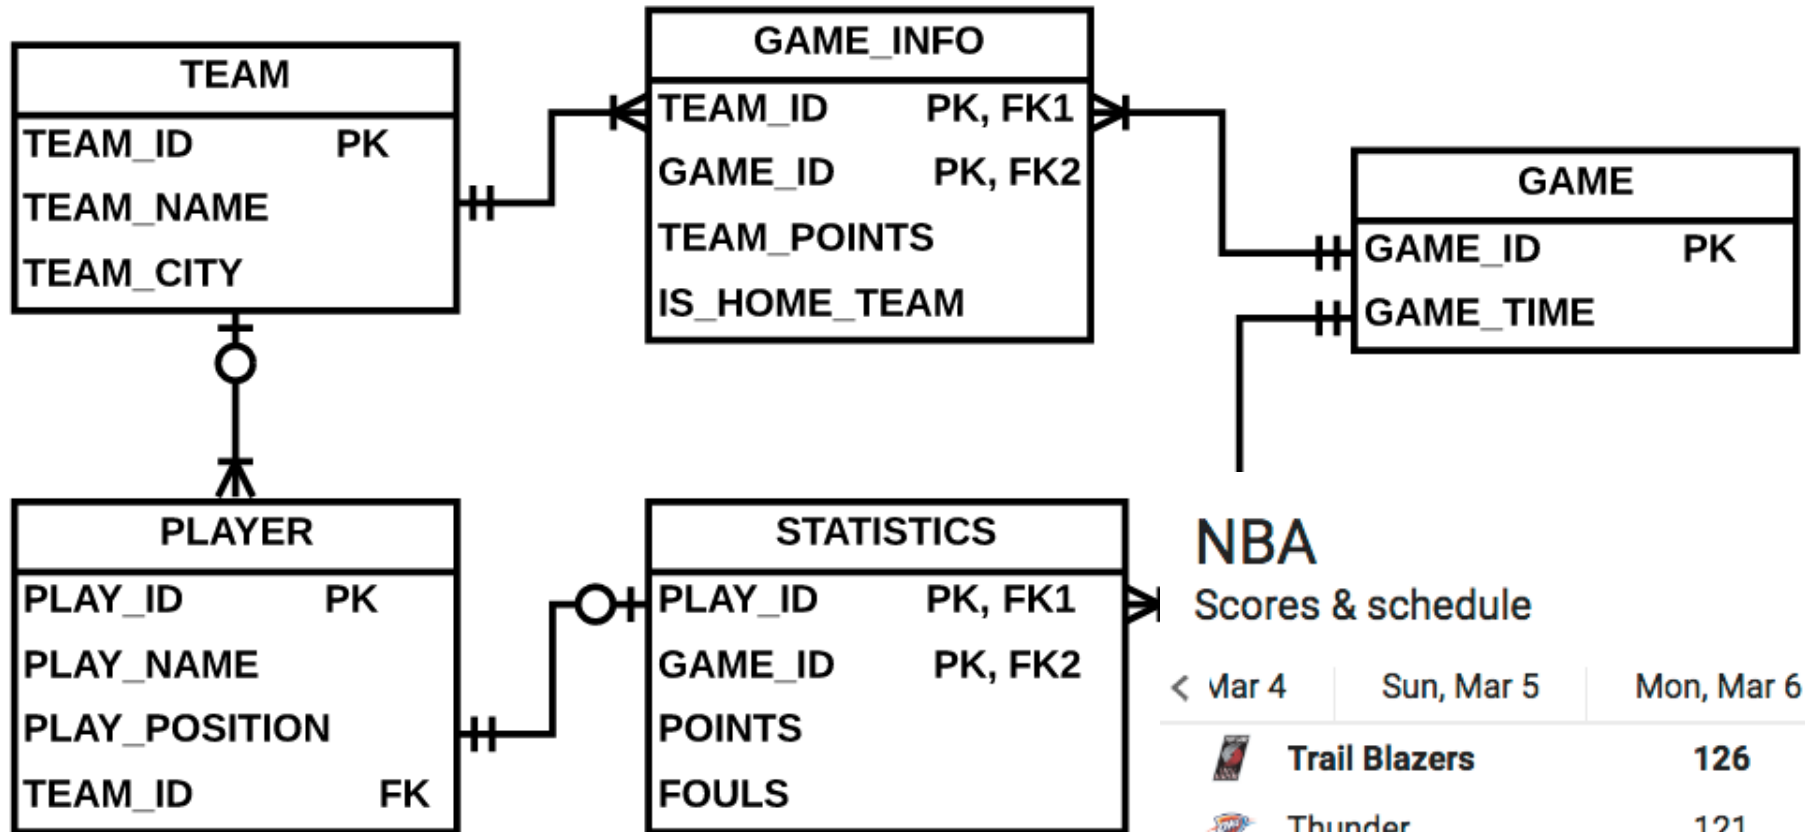
- Or use subquery

```
SELECT * FROM PAINTER WHERE PTR_NUM =  
(SELECT PTR_NUM FROM PAINTING WHERE PTNG_TITLE = 'Dawn Thunder');
```

Motivation (cont.)

- Can you list *not only* painter information *but also* gallery information of the painting 'Dawn Thunder'.
- Things become more complicated when data are retrieved from multiple tables.

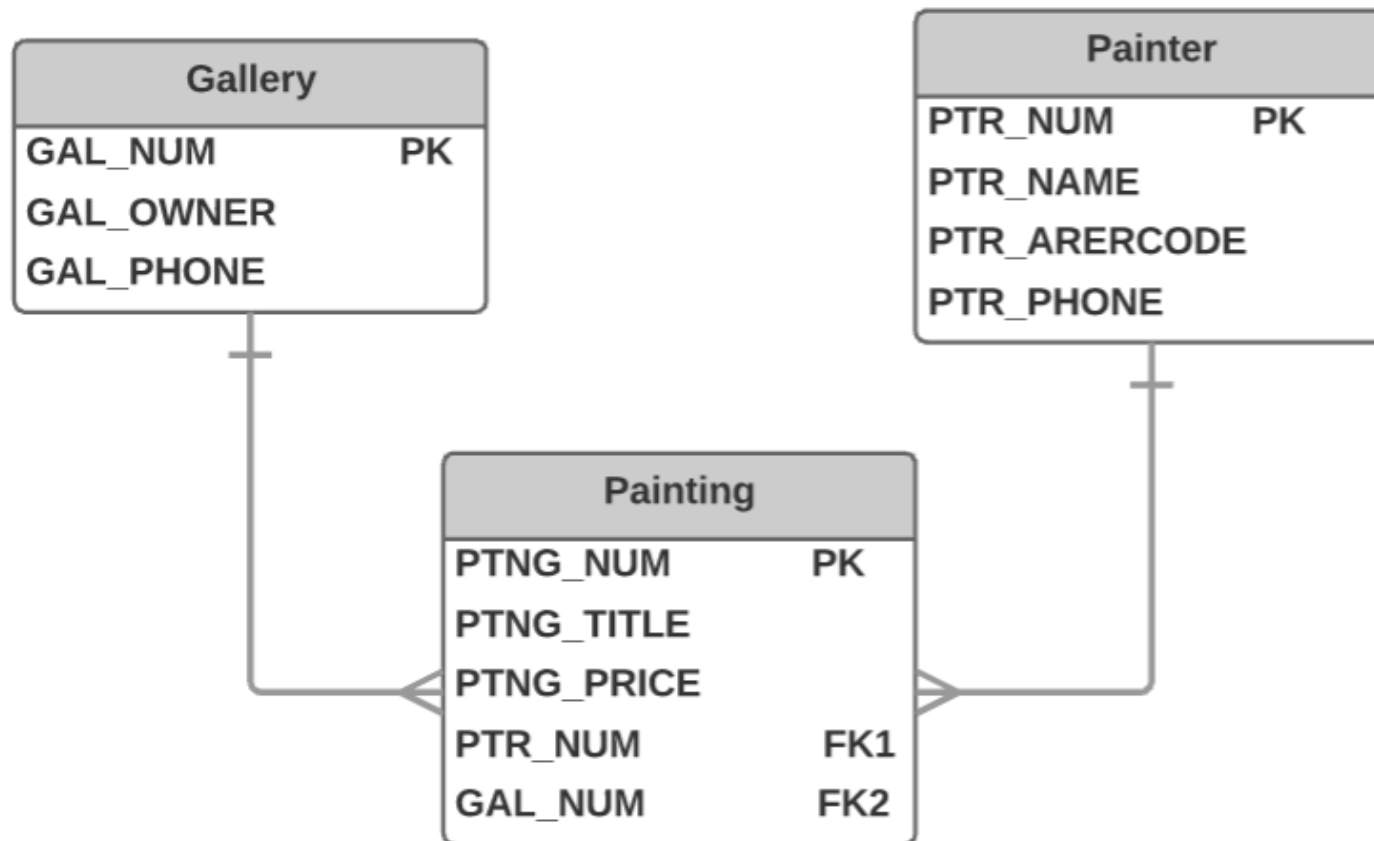
Motivation (cont.)



Query this database:

- How many games has Player A played?
- What's the average points that Player A has scored?

Recall PK, FK



PTNG_NUM	PTNG_TITLE	PTNG_PRICE	PTR_NUM	GAL_NUM
1338	Dawn Thunder	245.50	123	5
1339	A Faded Rose	6723.00	123	NULL
1340	The Founders	567.99	126	6
1341	Hasty Pudding Exit	145.50	123	NULL
1342	Plastic Paradise	8328.99	126	6
1343	Roamin'	785.00	127	6
1344	Wild Waters	999.00	127	5
1345	Stuff 'n Such 'n Some	9800.00	123	5

PAINTING

PTNG_NUM	PTNG_TITLE	PTNG_PRICE	PTR_NUM	PTR_NAME	PTR_AREACODE	PTR_PHONE	GAL_NUM	GAL_OWNER	GAL_PHONE
1338	Dawn Thunder	245.50	123	Ross	901	885-4567	5	Alice	123-4456
1339	A Faded Rose	6723.00	123	Ross	901	885-4567	NULL	NULL	NULL
1340	The Founders	567.99	126	Itero	901	346-1112	6	Waters	353-2243
1341	Hasty Pudding Exit	145.50	123	Ross	901	885-4567	NULL	NULL	NULL
1342	Plastic Paradise	8328.99	126	Itero	901	346-1112	6	Waters	353-2243
1343	Roamin'	785.00	127	Geoff	615	221-4456	6	Waters	353-2243
1344	Wild Waters	999.00	127	Geoff	615	221-4456	5	Alice	123-4456
1345	Stuff 'n Such 'n Some	9800.00	123	Ross	901	885-4567	5	Alice	123-4456

The importance of **Entity Integrity**, **Referential Integrity**!!

SQL Join

- A SQL **join** clause combines records from two or more tables in a relational database.
- Natural join, Inner join, left join, right join, outer join.

Natural Join

- A natural join returns all rows (of joined tables) with matching values in the matching columns (same column name and data type) and eliminates duplicate columns.
- Syntax:

```
SELECT columns FROM table1 NATURAL JOIN table2;
```

PTR_NUM	PTR_NAME	PTR_AREACODE	PTR_PHONE
123	Ross	901	885-4567
126	Itero	901	346-1112
127	Geoff	615	221-4456

PTNG_NUM	PTNG_TITLE	PTNG_PRICE	PTR_NUM	GAL_NUM
1338	Dawn Thunder	245.50	123	5
1339	A Faded Rose	6723.00	123	NULL
1340	The Founders	567.99	126	6
1341	Hasty Pudding Exit	145.50	123	NULL
1342	Plastic Paradise	8328.99	126	6
1343	Roamin'	785.00	127	6
1344	Wild Waters	999.00	127	5
1345	Stuff 'n Such 'n Some	9800.00	123	5

SELECT * FROM PAINTER NATURAL JOIN PAINTING;

PTR_NUM	PTNG_NUM	PTNG_TITLE	PTNG_PRICE	GAL_NUM	PTR_NAME	PTR_AREACODE	PTR_PHONE
123	1338	Dawn Thunder	245.50	5	Ross	901	885-4567
123	1339	A Faded Rose	6723.00	NULL	Ross	901	885-4567
126	1340	The Founders	567.99	6	Itero	901	346-1112
123	1341	Hasty Pudding Exit	145.50	NULL	Ross	901	885-4567
126	1342	Plastic Paradise	8328.99	6	Itero	901	346-1112
127	1343	Roamin'	785.00	6	Geoff	615	221-4456
127	1344	Wild Waters	999.00	5	Geoff	615	221-4456
123	1345	Stuff 'n Such 'n Some	9800.00	5	Ross	901	885-4567

Only one PTR_NUM column

GAL_NUM	GAL_OWNER	GAL_PHONE
5	Alice	123-4456
6	Waters	353-2243

PTNG_NUM	PTNG_TITLE	PTNG_PRICE	PTR_NUM	GAL_NUM
1338	Dawn Thunder	245.50	123	5
1339	A Faded Rose	6723.00	123	NULL
1340	The Founders	567.99	126	6
1341	Hasty Pudding Exit	145.50	123	NULL
1342	Plastic Paradise	8328.99	126	6
1343	Roamin'	785.00	127	6
1344	Wild Waters	999.00	127	5
1345	Stuff 'n Such 'n Some	9800.00	123	5

SELECT * FROM GALLERY NATURAL JOIN PAINTING;

GAL_NUM	GAL_OWNER	GAL_PHONE	PTNG_NUM	PTNG_TITLE	PTNG_PRICE	PTR_NUM
5	Alice	123-4456	1338	Dawn Thunder	245.50	123
6	Waters	353-2243	1340	The Founders	567.99	126
6	Waters	353-2243	1342	Plastic Paradise	8328.99	126
6	Waters	353-2243	1343	Roamin'	785.00	127
5	Alice	123-4456	1344	Wild Waters	999.00	127
5	Alice	123-4456	1345	Stuff 'n Such 'n Some	9800.00	123

NULL gallery is not in the return result

GAL_NUM	GAL_OWNER	GAL_PHONE
5	Alice	123-4456
6	Waters	353-2243

PTR_NUM	PTR_NAME	PTR_AREACODE	PTR_PHONE
123	Ross	901	885-4567
126	Itero	901	346-1112
127	Geoff	615	221-4456

SELECT * FROM GALLERY NATURAL JOIN PAINTER;

What if there is no matching column?
Return a **relational product** (cartesian product) of two tables.

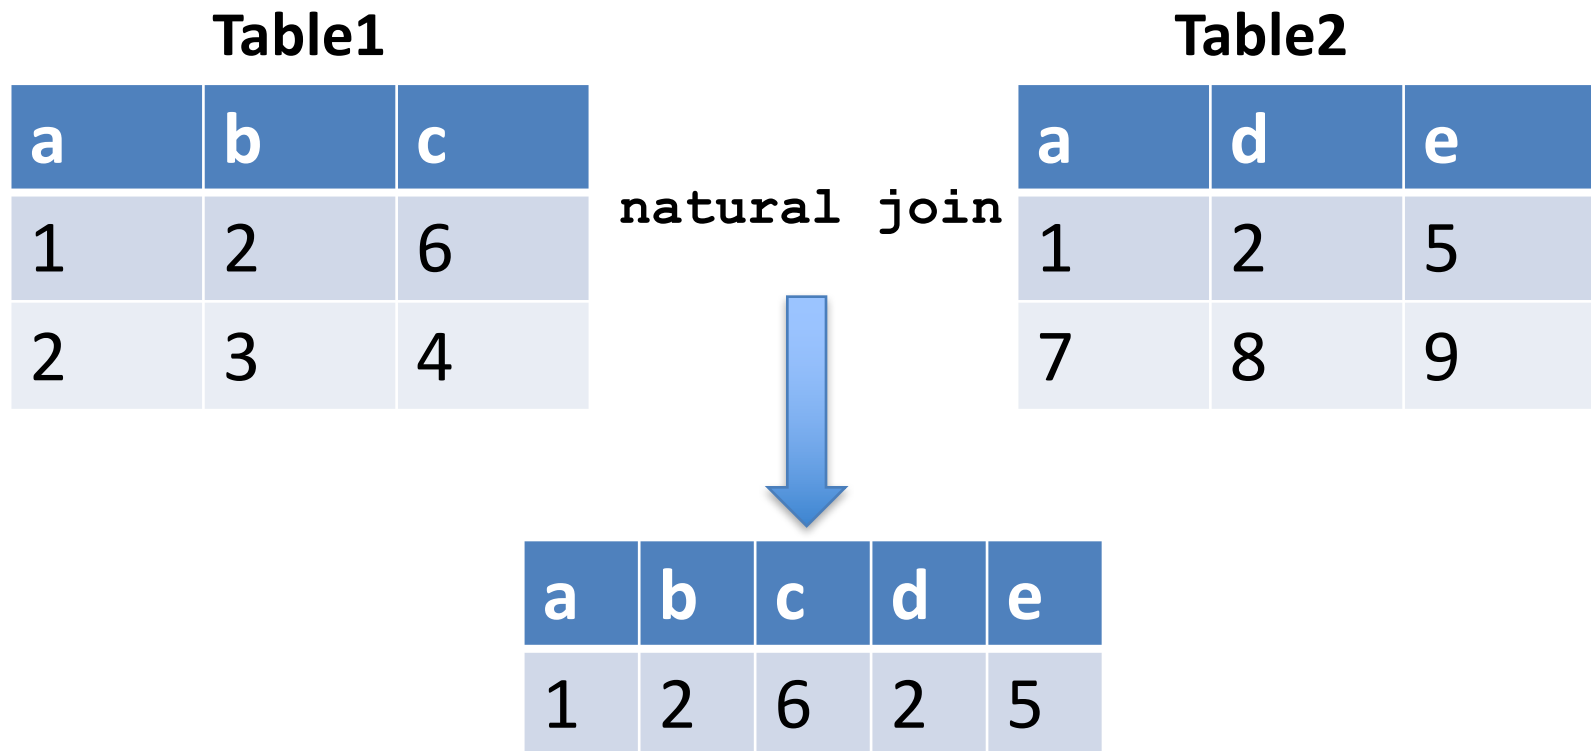
GAL_NUM	GAL_OWNER	GAL_PHONE	PTR_NUM	PTR_NAME	PTR_AREACODE	PTR_PHONE
5	Alice	123-4456	123	Ross	901	885-4567
6	Waters	353-2243	123	Ross	901	885-4567
5	Alice	123-4456	126	Itero	901	346-1112
6	Waters	353-2243	126	Itero	901	346-1112
5	Alice	123-4456	127	Geoff	615	221-4456
6	Waters	353-2243	127	Geoff	615	221-4456

Natural Join

- Natural join performs the following tasks
 - Determine the **common column(s)** by looking for columns with identical names and compatible data types;
 - Select only the rows with **common values** in the common columns.
 - If there are no common column, return the **relational product** of the tables.

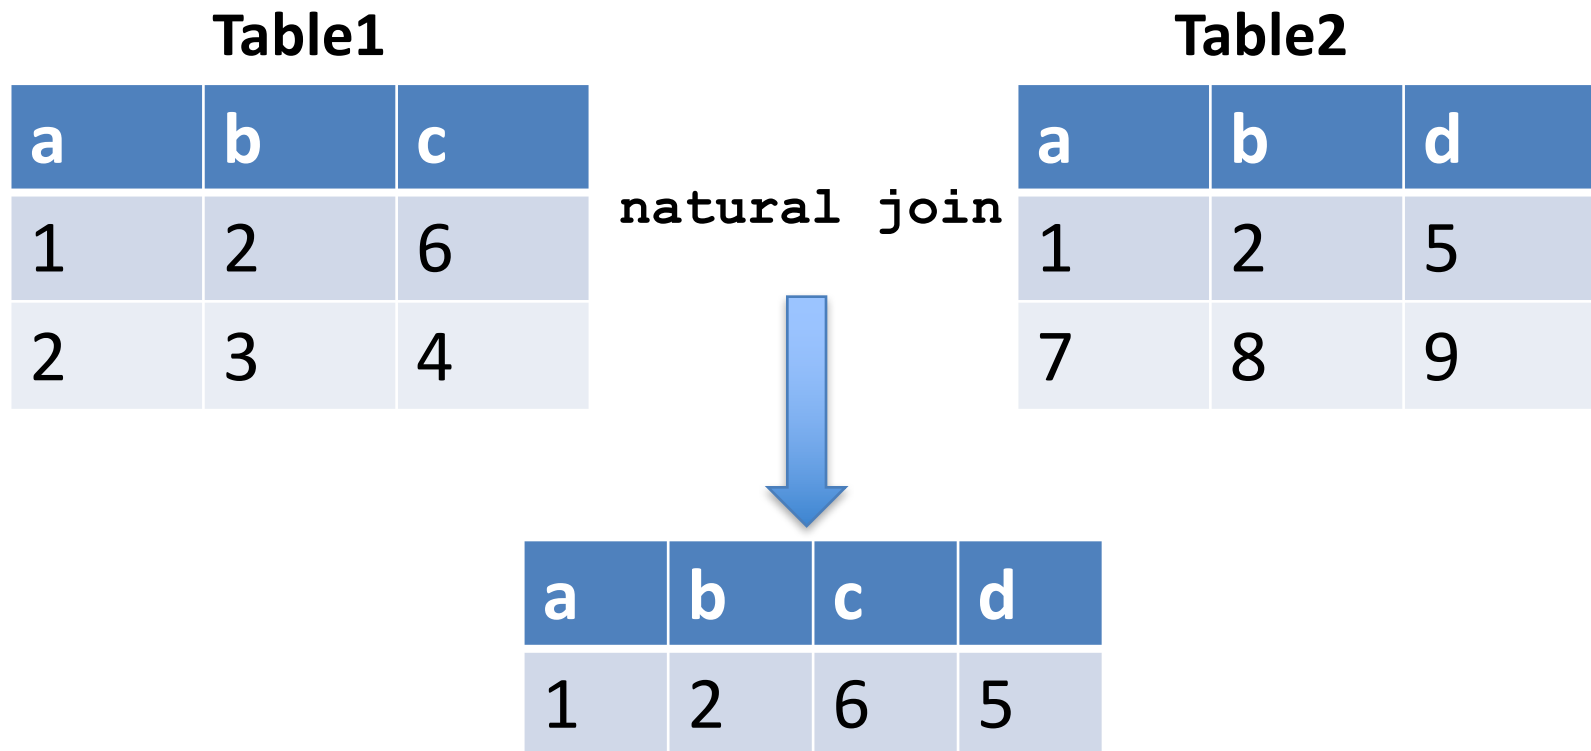
Natural Join

- One common column.



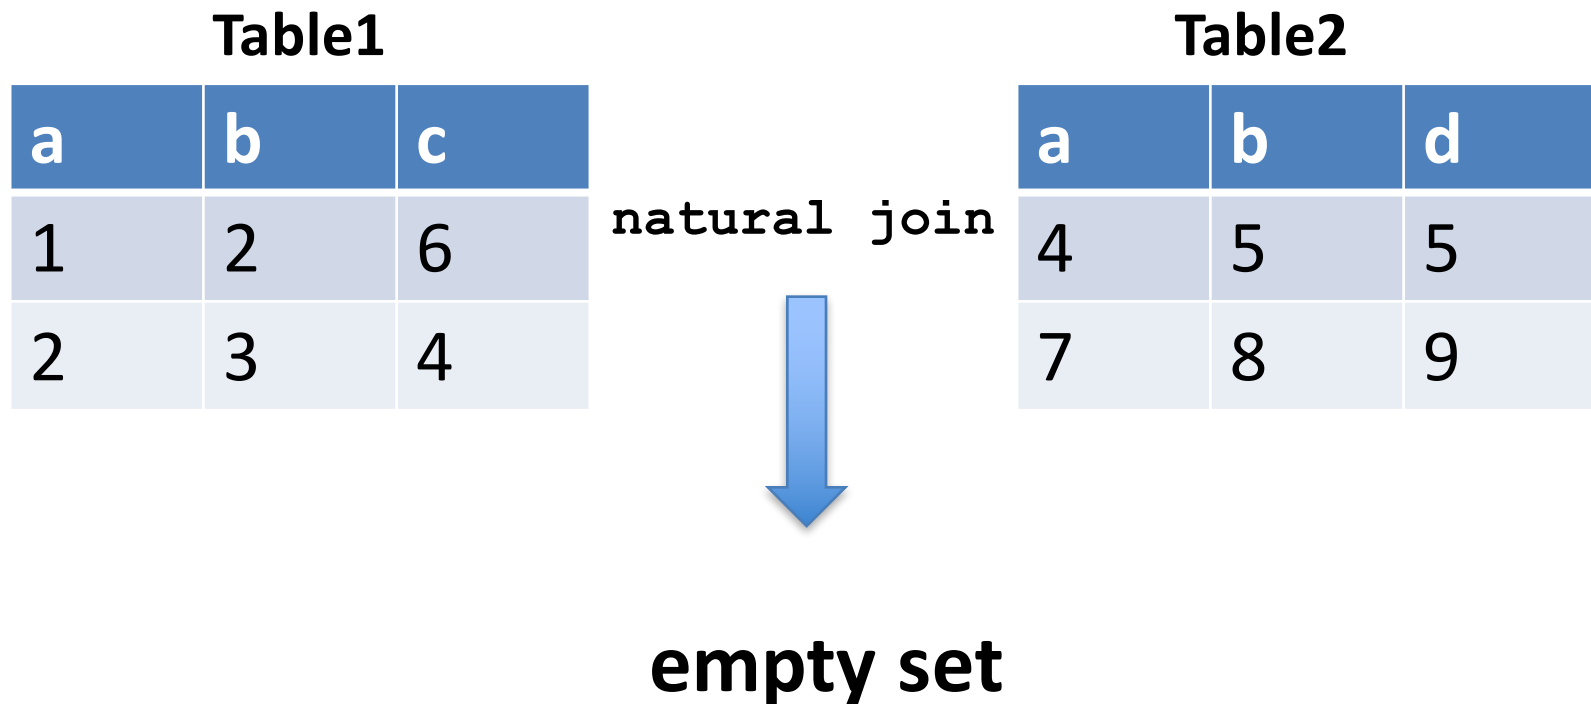
Natural Join (cont.)

- What if there are more than one common column?



Natural Join (cont.)

- What if there are no common values in the common columns?



Natural Join (cont.)

- What if there is no common columns?

Table1

a	b
1	2
2	3

natural join



Table2

c	d
4	5
7	8

a	b	c	d
1	2	4	5
1	2	7	8
2	3	4	5
2	3	7	8

Relational product

iClicker question 1

- How many rows will the following natural join return?

Table1

a	b
1	2
2	3

natural join

Table2

c	d
4	5
7	8
9	1

- a. 4
- b. 0
- c. 9
- d. 6

iClicker question 2

- How many rows will the following natural join return?

Table1

a	b
1	2
2	3

natural join

Table2

a	c
4	5
7	8
9	1

- a. 4
- b. 0
- c. 9
- d. 6

Structure Query Language

```
SELECT columns
FROM table
[WHERE condition]
[GROUP BY column]
[HAVING condition]
[ORDER BY column [ASC | DESC]]
[LIMIT count]
```

- Multiple tables: Natural Join Painter, Painting, Gallery.

```
SELECT * FROM PAINTING NATURAL JOIN PAINTER
NATURAL JOIN GALLERY;
```

- Where, GROUP BY, ORDER BY, etc...

```
SELECT COUNT(*), PAINTING.PTR_NUM FROM
PAINTING NATURAL JOIN PAINTER
GROUP BY PTR_NUM HAVING COUNT(*) > 2;
```

Derived table

```
SELECT columns FROM table  
[WHERE ....]  
[GROUP BY ....]  
[ORDER BY ....]
```

a real stored table

a derived table

```
SELECT columns FROM table3  
[WHERE ....]  
[GROUP BY ....]
```

```
SELECT MIN(P_PRICE) FROM  
(SELECT * FROM PRODUCT  
WHERE P_PRICE > (SELECT AVG(P_PRICE) FROM PRODUCT)) AS T;
```

- Although join clause generates a “virtual” table, it’s *not* a derived table because it’s not a subquery.

INNER JOIN

- What if two columns in two tables do not have the same name, but indeed they are the same thing, or
- what if the columns you don't want to join happen to have the same name?

UserID	Name	Email
A00001	Alice	alice@
A00002	Bob	bob@
A00003	Doug	doug1@
A00004	Doug	doug2@

Facebook

UserID	Name	EmailAdd
1	Alice	alice@
2	Doug	doug2@
3	Emma	emma@
4	Fox	fox@

Instagram

INNER JOIN (cont.)

- An inner join returns all rows (of two tables) with matching values in the columns that are specified in the condition.
- ```
SELECT columns FROM table1
INNER JOIN table2 ON join_condition;
```

Table qualifier: use a dot to refer a table

```
SELECT * FROM PAINTER INNER JOIN PAINTING
ON PAINTER.PTR_NUM = PAINTING.PTR_NUM;
```

| PTR_NUM | PTR_NAME | PTR_AREACODE | PTR_PHONE | PTNG_NUM | PTNG_TITLE            | PTNG_PRICE | PTR_NUM | GAL_NUM |
|---------|----------|--------------|-----------|----------|-----------------------|------------|---------|---------|
| 123     | Ross     | 901          | 885-4567  | 1338     | Dawn Thunder          | 245.50     | 123     | 5       |
| 123     | Ross     | 901          | 885-4567  | 1339     | A Faded Rose          | 6723.00    | 123     | NULL    |
| 126     | Itero    | 901          | 346-1112  | 1340     | The Founders          | 567.99     | 126     | 6       |
| 123     | Ross     | 901          | 885-4567  | 1341     | Hasty Pudding Exit    | 145.50     | 123     | NULL    |
| 126     | Itero    | 901          | 346-1112  | 1342     | Plastic Paradise      | 8328.99    | 126     | 6       |
| 127     | Geoff    | 615          | 221-4456  | 1343     | Roamin'               | 785.00     | 127     | 6       |
| 127     | Geoff    | 615          | 221-4456  | 1344     | Wild Waters           | 999.00     | 127     | 5       |
| 123     | Ross     | 901          | 885-4567  | 1345     | Stuff 'n Such 'n Some | 9800.00    | 123     | 5       |

Natural Join and Inner join with the same condition will return exactly the same rows, but more columns.

# table qualifier

```
SELECT * FROM PAINTER INNER JOIN PAINTING
ON PAINTER.PTR_NUM = PAINTING.PTR_NUM;
```

# table alias

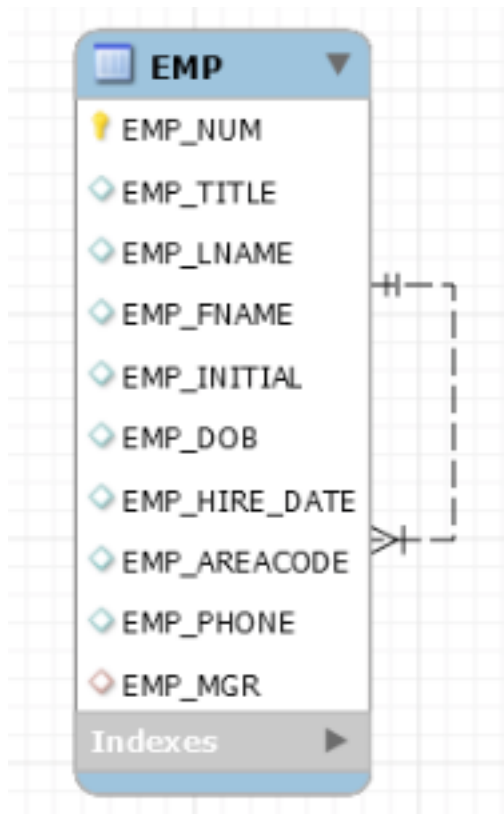
```
SELECT * FROM PAINTER AS P INNER JOIN PAINTING AS Q
ON P.PTR_NUM = Q.PTR_NUM;
```



Table alias

# Table join itself

List all employees with their managers' names



| EMP_NUM | EMP_LNAME  | EMP_FNAME | EMP_MGR |
|---------|------------|-----------|---------|
| 100     | Kolmycz    | George    | NULL    |
| 102     | VanDam     | Rhett     | 100     |
| 103     | Jones      | Anne      | 100     |
| 104     | Lange      | John      | 105     |
| 105     | Williams   | Robert    | NULL    |
| 106     | Smith      | Jeanine   | 105     |
| 107     | Diante     | Jorge     | 105     |
| 108     | Wiesenbach | Paul      | NULL    |
| 109     | Smith      | George    | 108     |
| 110     | Genkazi    | Leighla   | 108     |
| 111     | Washington | Rupert    | 105     |
| 112     | Johnson    | Edward    | 100     |
| 113     | Smythe     | Melanie   | 105     |
| 114     | Brandon    | Marie     | 108     |
| 115     | Saranda    | Hermine   | 105     |
| 116     | Smith      | George    | 108     |
| 101     | Lewis      | Rhonda    | 100     |

# Table join itself

| EMP_NUM | EMP_LNAME  | EMP_FNAME | EMP_NUM | EMP_LNAME  | EMP_FNAME |
|---------|------------|-----------|---------|------------|-----------|
| 102     | VanDam     | Rhett     | 100     | Kolmycz    | George    |
| 103     | Jones      | Anne      | 100     | Kolmycz    | George    |
| 112     | Johnson    | Edward    | 100     | Kolmycz    | George    |
| 101     | Lewis      | Rhonda    | 100     | Kolmycz    | George    |
| 104     | Lange      | John      | 105     | Williams   | Robert    |
| 106     | Smith      | Jeanine   | 105     | Williams   | Robert    |
| 107     | Diante     | Jorge     | 105     | Williams   | Robert    |
| 111     | Washington | Rupert    | 105     | Williams   | Robert    |
| 113     | Smythe     | Melanie   | 105     | Williams   | Robert    |
| 115     | Saranda    | Hermine   | 105     | Williams   | Robert    |
| 109     | Smith      | George    | 108     | Wiesenbach | Paul      |
| 110     | Genkazi    | Leighla   | 108     | Wiesenbach | Paul      |
| 114     | Brandon    | Marie     | 108     | Wiesenbach | Paul      |
| 116     | Smith      | George    | 108     | Wiesenbach | Paul      |

```
SELECT M.EMP_NUM, M.EMP_LNAME, M.EMP_FNAME,
N.EMP_NUM, N.EMP_LNAME, N.EMP_FNAME
FROM EMP M INNER JOIN EMP N ON M.EMP_MGR = N.EMP_NUM;
```

- Alias is needed when a table is joined to itself.

# Difference between NATURAL JOIN and INNER JOIN

- “INNER JOIN is often used in real life.”
- NATRUAL JOIN is an inner join involving all columns that happen to have the same name.
- They return different number of columns.  
(Natural join eliminates duplicated columns.)

# Summary

- **JOIN**, a very powerful command in relational database.
- Natural Join
- Inner Join